

Community ecology- Computer lab I - AB332

Ramiro Logares (ICM), with additions by Anders K. Krabberød (UiO)

October 2021

This is an [R Markdown](#) Notebook. When you execute code within the notebook, the results appear beneath the code. Try executing this chunk by clicking the *Run* button within the chunk or by placing your cursor inside it and pressing *Cmd+Shift+Enter*.

Add a new chunk by clicking the *Insert Chunk* button on the toolbar or by pressing *Cmd+Option+I*.

When you save the notebook, an HTML file containing the code and output will be saved alongside it (click the *Preview* button or press *Cmd+Shift+K* to preview the HTML file).

The preview shows you a rendered HTML copy of the contents of the editor.

Consequently, unlike *Knit*, *Preview* does not run any R code chunks. Instead, the output of the chunk when it was last run in the editor is displayed.

It's time for you to try and do the same analysis as was shown in the lecture but using a different dataset: Make sure you have installed all packages!

Load Packages

Starting community ecology analyses

Read the data from the github page:

```
otu.tab<-read_tsv("https://raw.githubusercontent.com/krabberod/UNIS_AB332_2021/main/computer_lab/data/AB332_otutab_reduc3.txt")

## Rows: 3697 Columns: 83

## -- Column specification -----
## Delimiter: "\t"
## chr (1): OTUNumber
## dbl (82): Isa_111214, Isa_120117, Isa_120128, Isa_120209, Isa_120216, Isa_12...

##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

First, get to know the data: How many samples and how many OTUs are in the dataset? What do the numbers in the sample names mean?

```
head(otu.tab)

## # A tibble: 6 x 83
##   OTUNumber Isa_111214 Isa_120117 Isa_120128 Isa_120209 Isa_120216 Isa_120223
##   <chr>         <dbl>     <dbl>     <dbl>     <dbl>     <dbl>     <dbl>
## 1 OTU1          436       348       236       139       260       446
## 2 OTU10        2308      2537      1599      1956      328       668
```

```
## 3 OTU100      83      36      21      30      1      21
## 4 OTU1000     0       0       0       0       0       0
## 5 OTU1001     1       0       0       0       0       2
## 6 OTU1002     0       2       0       1       0       0
## # ... with 76 more variables: Isa_120301 <dbl>, Isa_120308 <dbl>,
## #   Isa_120320 <dbl>, Isa_120321 <dbl>, Isa_120322 <dbl>, Isa_120323 <dbl>,
## #   Isa_120329 <dbl>, Isa_120403 <dbl>, Isa_120411 <dbl>, Isa_120416 <dbl>,
## #   Isa_120419 <dbl>, Isa_120423 <dbl>, Isa_120426 <dbl>, Isa_120430 <dbl>,
## #   Isa_120503 <dbl>, Isa_120507 <dbl>, Isa_120508 <dbl>, Isa_120509 <dbl>,
## #   Isa_120510 <dbl>, Isa_120516 <dbl>, Isa_120524 <dbl>, Isa_120621 <dbl>,
## #   Isa_120706 <dbl>, Isa_120806 <dbl>, Isa_120823 <dbl>, Isa_120906 <dbl>, ...

dim(otu.tab)

## [1] 3697  83
```

You can look at a given selection of the table by specifying a range of rows and columns:

```
otu.tab[5:15,1:5] # The first 10 rows, and the first 5 columns

## # A tibble: 11 x 5
##   OTUNumber Isa_111214 Isa_120117 Isa_120128 Isa_120209
##   <chr>      <dbl>      <dbl>      <dbl>      <dbl>
## 1 OTU1001      1         0         0         0
## 2 OTU1002      0         2         0         1
## 3 OTU1003      3         0         0         0
## 4 OTU1004      9         2         5         7
## 5 OTU1005      1         0         2         0
## 6 OTU1006      0         2         0         4
## 7 OTU1007      0         0         0         0
## 8 OTU1008      0         0         0         0
## 9 OTU1009      0         0         0         1
## 10 OTU101      3         0         0         0
## 11 OTU1010     1         3         0         3
```

You can also see the entire table with the View() function:

```
View(otu.tab)
```

See if you can choose a different subset. For instance samples 6-12 and Otus 20-26:

We assign OTU numbers as rownames

```
otu.tab <- column_to_rownames(otu.tab, var = "OTUNumber")
```

Let's check the names

```
head(rownames(otu.tab))

## [1] "OTU1" "OTU10" "OTU100" "OTU1000" "OTU1001" "OTU1002"

dim(otu.tab)

## [1] 3697  82
```

For simplicity, I have included only the 25 samples in the rest of the tutorial. As an exercise, you should redo the analysis with the full dataset. I.e. remove the part of the code that selects samples

1:15 in the following chunk. (This way your numbers will differ from the pdf, and you can also see the effect of a different dataset).

```
otu.tab.red<-otu.tab[,6:30]
```

The data needs to be transposed since this is how Vegan likes it.

```
otu.tab.simple<-t(otu.tab.red)
otu.tab.simple[1:5,1:5]
```

```
##           OTU1 OTU10 OTU100 OTU1000 OTU1001
## Isa_120223  446   668     21        0        2
## Isa_120301  149   551     23        1        1
## Isa_120308  321  1462     99       11        0
## Isa_120320  204   896     75        1        0
## Isa_120321  442   646     61        4        0
```

You can get the total number of reads for each sample using rowSums(), and the total reads per OTU with colSums()

```
rowSums(otu.tab.simple)
```

```
## Isa_120223 Isa_120301 Isa_120308 Isa_120320 Isa_120321 Isa_120322 Isa_120323
##      13682      15783      32833      19361      17110      13658      16251
## Isa_120329 Isa_120403 Isa_120411 Isa_120416 Isa_120419 Isa_120423 Isa_120426
##      17002      13551      25606      29877      16194      17161      22305
## Isa_120430 Isa_120503 Isa_120507 Isa_120508 Isa_120509 Isa_120510 Isa_120516
##      19942      25237      27389      19969      34666      17682      25256
## Isa_120524 Isa_120621 Isa_120706 Isa_120806
##      40267      24636      23156      20390
```

```
head(colSums(otu.tab.simple)) # Too many to show them all.
```

```
##      OTU1      OTU10      OTU100 OTU1000 OTU1001 OTU1002
## 100420      7814       786        54        4        28
```

Since I have selected only a few of the samples it is possible that some of the OTU's are left with a total abundance of zero. In R it is possible to have functions within functions so the following will print the number of columns in the data set that has a sum equal to 0:

```
length(which(colSums(otu.tab.simple)==0))
```

```
## [1] 1163
```

We can use the same idea of a function within a function to exclude the OTUs with a total number of 0.

```
otu.tab.simple<-otu.tab.simple[,-(which(colSums(otu.tab.simple)==0))]
```

Now how many are 0?

```
length(which(colSums(otu.tab.simple)==0))
```

```
## [1] 0
```

How many have more than 0 reads?

```
length(which(colSums(otu.tab.simple)>0))

## [1] 2534
```

Can you find how many OTU's that have more than 10 reads (in total)?

Common metrics and methods

The following calculations make use of functions in the vegan package written by Jari Oksanen. *Vegan is an R package for community ecologists. It contains the most popular methods of multivariate analysis needed in analysing ecological communities, and tools for diversity analysis, and other potentially useful functions.* If you want to learn more about the vegan you can run `browseVignettes("vegan")`

Richness estimations

```
richness<-estimateR(otu.tab.simple)
richness
```

##	Isa_120223	Isa_120301	Isa_120308	Isa_120320	Isa_120321	Isa_120322
## S.obs	701.00000	902.00000	1042.00000	952.00000	897.00000	858.00000
## S.chao1	1100.22581	1260.69343	1355.30612	1316.81633	1200.76642	1337.40517
## se.chao1	66.62128	54.17742	47.66939	53.77784	47.47591	72.14089
## S.ACE	1036.24030	1249.23447	1337.63473	1341.50902	1197.73116	1278.33686
## se.ACE	17.06514	18.59112	18.50405	19.72472	17.80976	19.74198
##	Isa_120323	Isa_120329	Isa_120403	Isa_120411	Isa_120416	Isa_120419
## S.obs	959.00000	942.00000	848.00000	1088.00000	992.00000	450.00000
## S.chao1	1377.95000	1397.55469	1155.65000	1517.36486	1287.23077	731.44286
## se.chao1	60.96351	67.07142	47.65508	61.13897	44.73270	54.74508
## S.ACE	1365.23150	1356.87666	1181.99016	1480.20490	1295.47022	759.36594
## se.ACE	19.72705	19.88369	18.05970	20.02279	18.55676	16.47752
##	Isa_120423	Isa_120426	Isa_120430	Isa_120503	Isa_120507	Isa_120508
## S.obs	384.00000	374.00000	304.00000	420.00000	272.00000	219.00000
## S.chao1	693.25532	511.60000	438.63830	654.23077	447.77778	324.63636
## se.chao1	67.46164	32.24236	34.05155	48.23404	45.97426	31.16544
## S.ACE	652.06713	539.26272	435.54374	684.48661	432.26064	329.85163
## se.ACE	14.78726	12.98133	10.86998	15.37044	11.73649	9.56906
##	Isa_120509	Isa_120510	Isa_120516	Isa_120524	Isa_120621	Isa_120706
## S.obs	205.00000	224.00000	252.00000	335.00000	417.00000	414.00000
## S.chao1	363.05263	371.17143	369.00000	528.75000	618.88235	582.52174
## se.chao1	51.97189	40.07574	33.24605	48.16268	46.10172	36.35228
## S.ACE	304.84401	395.87277	364.99135	514.52628	590.13947	615.28756
## se.ACE	9.30590	12.00009	10.37553	12.39624	13.01145	13.57143
##	Isa_120806					
## S.obs	605.00000					
## S.chao1	873.07692					
## se.chao1	50.85980					
## S.ACE	851.28562					
## se.ACE	15.83698					

Above we have the estimators Chao and ACE as well as the species number. What do the numbers mean?

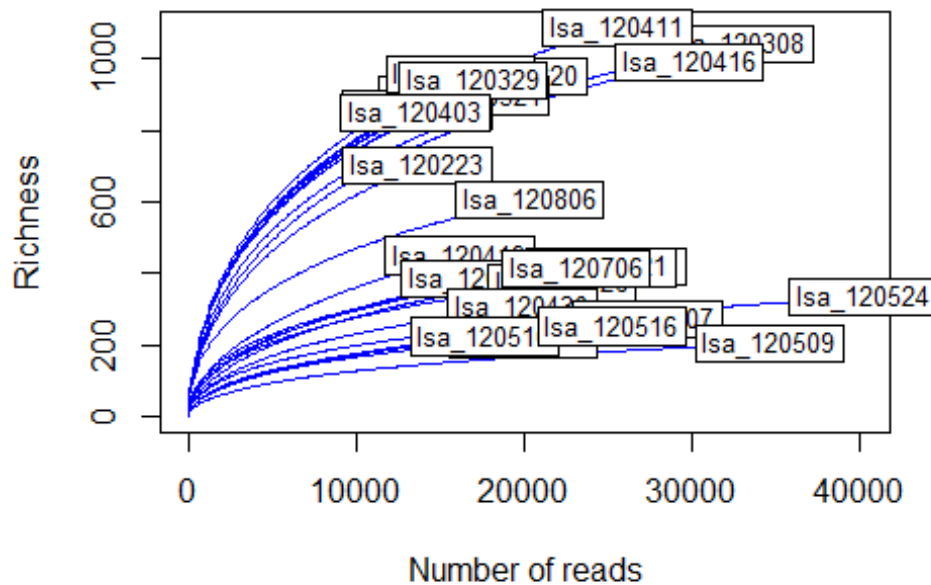
Rarefaction

Let's calculate the number of reads per sample.

```
rowSums(otu.tab.simple)
```

```
## Isa_120223 Isa_120301 Isa_120308 Isa_120320 Isa_120321 Isa_120322 Isa_120323
##      13682      15783      32833      19361      17110      13658      16251
## Isa_120329 Isa_120403 Isa_120411 Isa_120416 Isa_120419 Isa_120423 Isa_120426
##      17002      13551      25606      29877      16194      17161      22305
## Isa_120430 Isa_120503 Isa_120507 Isa_120508 Isa_120509 Isa_120510 Isa_120516
##      19942      25237      27389      19969      34666      17682      25256
## Isa_120524 Isa_120621 Isa_120706 Isa_120806
##      40267      24636      23156      20390
```


```
rarecurve (otu.tab.simple, step=100, xlab= "Number of reads", ylab="Richness", col="blue")
```



How do you interpret these curves? Which samples have the lowest number of total reads? Which are the highest?

Accumulation curves

```
accum.curve<-specaccum(otu.tab.simple, method="collector")
plot(accum.curve)
```

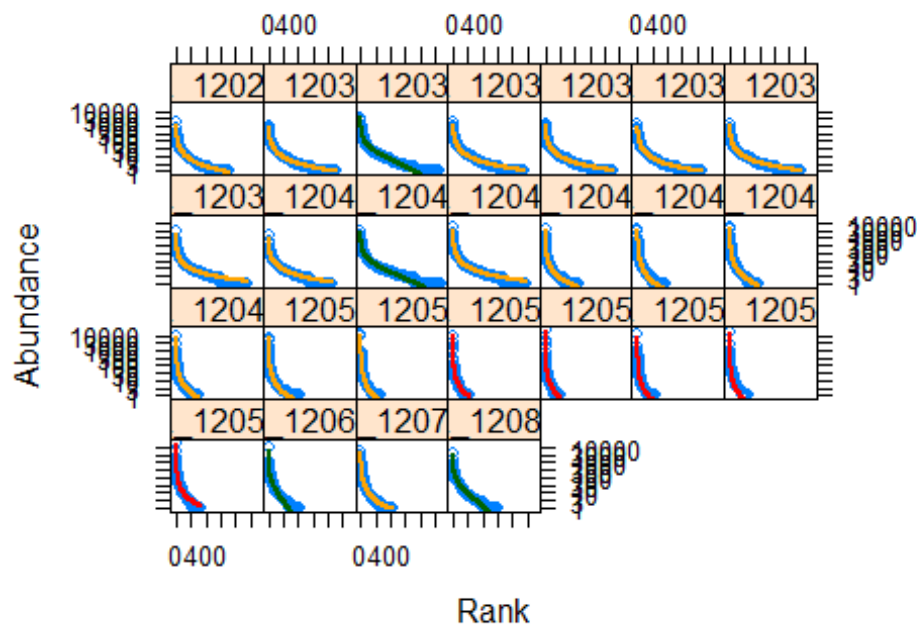


What does this curve represent? How do you interpret it?

Evenness

```
plot(colSums(otu.tab.simple), log="y", xlab="Rank", ylab="Abundance", pch=19, cex=0.5, col="blue")
```


lull ————— Lognormal ————— Mandelbrot —————
 'reemption ————— Zipf —————

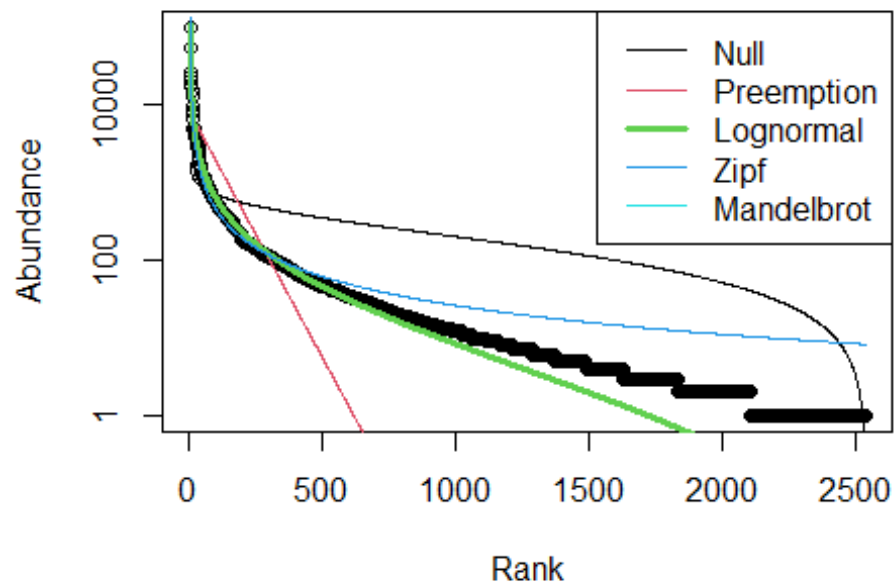


```

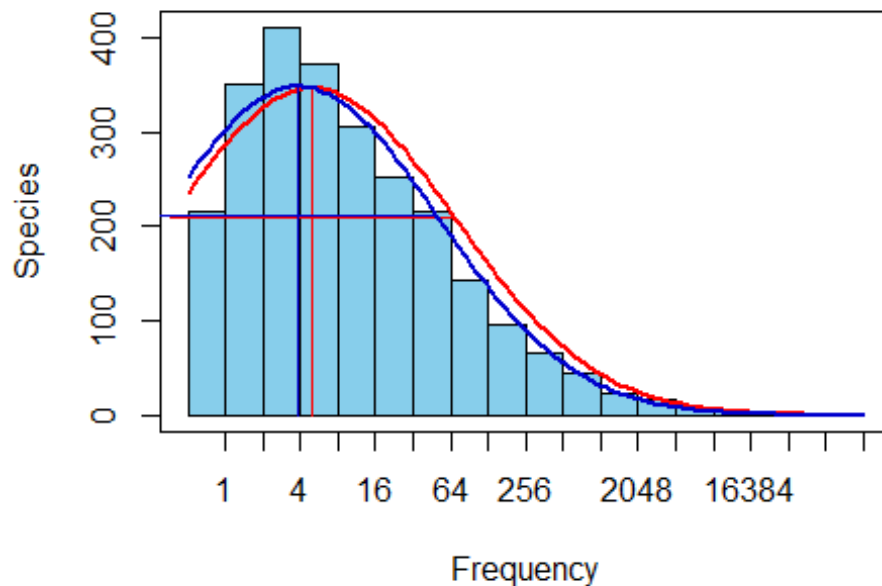
mod.all<-radfit(colSums(otu.tab.simple))

## Error in glm.fit(x = structure(c(1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  :
##   NA/NaN/Inf in 'x'

plot(mod.all)
  
```

```
#Fitting data to the Preston model
preston<-prestonfit(colSums(otu.tab.simple))
preston.dist<-prestondistr(colSums(otu.tab.simple))
plot(preston)
lines(preston.dist, line.col="blue3")
```



Extrapolated richness

```
veiledspec(preston)
```

```
## Extrapolated      Observed      Veiled
##      3278.6837      2534.0000      744.6837
```

```
veiledspec(veiledspec.preston)
```

```
## Extrapolated      Observed      Veiled
##      3211.9212      2534.0000      677.9212
```

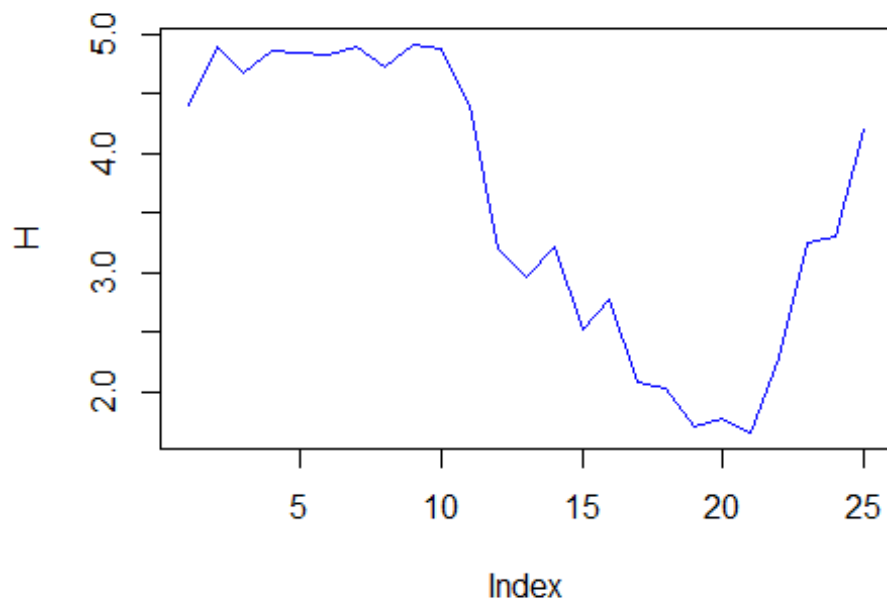
Shannon H index (considers richness and evenness)

```
H<-diversity(otu.tab.simple, index="shannon")
```

```
H
```

```
## Isa_120223 Isa_120301 Isa_120308 Isa_120320 Isa_120321 Isa_120322 Isa_120323
##      4.407409      4.902868      4.679546      4.859799      4.844499      4.826550      4.900137
## Isa_120329 Isa_120403 Isa_120411 Isa_120416 Isa_120419 Isa_120423 Isa_120426
##      4.725430      4.914670      4.872817      4.400562      3.203907      2.970289      3.211268
## Isa_120430 Isa_120503 Isa_120507 Isa_120508 Isa_120509 Isa_120510 Isa_120516
##      2.520753      2.781324      2.086866      2.028217      1.709035      1.784380      1.662585
## Isa_120524 Isa_120621 Isa_120706 Isa_120806
##      2.295421      3.245357      3.311757      4.199123
```

```
plot(H, type="l", col="blue")
```



Pielou's index of evenness (range 0-1, 1 = maximum evenness)

$J = H/H_{max}$


$J = \text{Shannon } (H) / \log(S = \text{species richness})$

```
J <- H/log(rowSums(otu.tab.simple>0))
```

Inverse Simpson's D index (richness+evenness. Larger values, larger diversity)

```
inv.simpson<-diversity(otu.tab.simple, "invsimpson")
```

```
plot(inv.simpson, type="l", col="blue")
```



Beta diversity

We rarefy all samples to the same sequencing depth, to reduce biases

```
min(rowSums(otu.tab.simple)) # We calculate the sample with the minimum amount of reads
## [1] 13551

otu.tab.simple.ss<-rrarefy(otu.tab.simple, min(rowSums(otu.tab.simple))) #Samples are rarefied to lowest number of reads
rowSums(otu.tab.simple.ss)

## Isa_120223 Isa_120301 Isa_120308 Isa_120320 Isa_120321 Isa_120322 Isa_120323
##      13551      13551      13551      13551      13551      13551      13551
## Isa_120329 Isa_120403 Isa_120411 Isa_120416 Isa_120419 Isa_120423 Isa_120426
##      13551      13551      13551      13551      13551      13551      13551
## Isa_120430 Isa_120503 Isa_120507 Isa_120508 Isa_120509 Isa_120510 Isa_120516
##      13551      13551      13551      13551      13551      13551      13551
## Isa_120524 Isa_120621 Isa_120706 Isa_120806
##      13551      13551      13551      13551
```

What is the number of reads these samples have been rarefied to? What does it imply, do you understand how it is done?

Check that the number of OTUs are the same in the new table

```
dim(otu.tab.simple)
```

```
## [1] 25 2534
dim(otu.tab.simple.ss)
## [1] 25 2534
```

The tables have the same size, but, after removing reads, several OTUs might be left with zero read abundance.

```
length(which(colSums(otu.tab.simple)==0))
## [1] 0
length(which(colSums(otu.tab.simple.ss)==0))
## [1] 198
head(which(colSums(otu.tab.simple.ss)==0)) # Show the OTUs and the position in the table that have 0 abundance for the first OTUs
## OTU1016 OTU1032 OTU1043 OTU1075 OTU1103 OTU1221
##      21      37      45      71     100     199
```

We can compare the number of reads for one of the OTUs:

```
colnames(otu.tab.simple)[13]
## [1] "OTU1009"
otu.tab.simple[,13] # This gives the abundance of the OTU1009 across the different samples in the table that is NOT subsampled
## Isa_120223 Isa_120301 Isa_120308 Isa_120320 Isa_120321 Isa_120322 Isa_120323
##      0      0      0      0      0      0      0
## Isa_120329 Isa_120403 Isa_120411 Isa_120416 Isa_120419 Isa_120423 Isa_120426
##      0      0      0      0      0      0      0
## Isa_120430 Isa_120503 Isa_120507 Isa_120508 Isa_120509 Isa_120510 Isa_120516
##      0      0      0      0      0      0      0
## Isa_120524 Isa_120621 Isa_120706 Isa_120806
##      0      0      0      3
otu.tab.simple.ss[,13] # This gives the abundance of the OTU1009 across the different samples in the table that IS subsampled
## Isa_120223 Isa_120301 Isa_120308 Isa_120320 Isa_120321 Isa_120322 Isa_120323
##      0      0      0      0      0      0      0
## Isa_120329 Isa_120403 Isa_120411 Isa_120416 Isa_120419 Isa_120423 Isa_120426
##      0      0      0      0      0      0      0
## Isa_120430 Isa_120503 Isa_120507 Isa_120508 Isa_120509 Isa_120510 Isa_120516
##      0      0      0      0      0      0      0
## Isa_120524 Isa_120621 Isa_120706 Isa_120806
##      0      0      0      3
```

We can remove the OTUs with zero abundance with a similar command as we used at the beginning of the lab:

```
otu.tab.simple.ss.nozero<-otu.tab.simple.ss[,-(which(colSums(otu.tab.simple.ss)==0))]  
# Removes OTUs with zero abundance  
length(which(colSums(otu.tab.simple.ss.nozero)==0)) # Check that no zero abundance OT  
Us are left  
## [1] 0
```

Let's check dimensions of the tables:

```
dim(otu.tab.simple.ss)  
## [1] 25 2534  
dim(otu.tab.simple.ss.nozero)  
## [1] 25 2336
```

2548-2226 = 322 , This is the number of OTUs that we expected to be removed. ## Compositional data analyses Replace zeros (problems with log calculations) with pseudo-counts

```
otu.tab.simple.gbm<-cmultRepl(t(otu.tab.simple), output = "p-counts")  
## No. corrected values: 48248  
otu.tab.simple.gbm[1:5,1:5] # We have a look to the replaced values  
##          Isa_120223 Isa_120301 Isa_120308 Isa_120320 Isa_120321  
## OTU1      4.460000e+02      149 3.210000e+02 2.040000e+02 4.420000e+02  
## OTU10     6.680000e+02      551 1.462000e+03 8.960000e+02 6.460000e+02  
## OTU100    2.100000e+01      23 9.900000e+01 7.500000e+01 6.100000e+01  
## OTU1000   7.334094e-03      1 1.100000e+01 1.000000e+00 4.000000e+00  
## OTU1001   2.000000e+00      1 1.024907e-03 1.230456e-03 3.803261e-04
```

centered log-ratio (clr) transformation

```
otu.tab.simple.gbm.clr<-clr(otu.tab.simple.gbm) # We apply a centered log-ratio (clr)  
transformation  
otu.tab.simple.gbm.clr[1:5,1:5] #Values now look different than counts.  
##          Isa_120223 Isa_120301 Isa_120308 Isa_120320 Isa_120321  
## OTU1      -1.084952 -2.181324 -1.413829 -1.8671505 -1.093961  
## OTU10     2.667074 2.474521 3.450347 2.9607266 2.633586  
## OTU100    1.549500 1.640472 3.100098 2.8224660 2.615852  
## OTU1000   -3.370798 1.544424 3.942319 1.5444236 2.930718  
## OTU1001   7.237552 6.544405 -0.338748 -0.1559649 -1.330076  
## attr(,"class")  
## [1] "rmult"
```

Distance metrics

First calculate the Bray Curtis dissimilarities for the rarefied dataset

```
otu.tab.simple.ss.nozero.bray<-vegdist(otu.tab.simple.ss.nozero, method="bray")  
as.matrix(otu.tab.simple.ss.nozero.bray)[1:5,1:5]  
##          Isa_120223 Isa_120301 Isa_120308 Isa_120320 Isa_120321  
## Isa_120223 0.0000000 0.3075788 0.2680245 0.2911224 0.2962143  
## Isa_120301 0.3075788 0.0000000 0.2427865 0.2030846 0.2073648
```

```
## Isa_120308 0.2680245 0.2427865 0.0000000 0.2611615 0.2900155
## Isa_120320 0.2911224 0.2030846 0.2611615 0.0000000 0.1329053
## Isa_120321 0.2962143 0.2073648 0.2900155 0.1329053 0.0000000
```

Then calculate the Euclidean distance based on the clr data (also known as Aitchison distance)

```
otu.tab.simple.gbm.clr.euclidean<-dist(t(otu.tab.simple.gbm.clr), method = "euclidean")
as.matrix(otu.tab.simple.gbm.clr.euclidean)[1:5,1:5]

##           Isa_120223 Isa_120301 Isa_120308 Isa_120320 Isa_120321
## Isa_120223      0.0000    172.3006    181.9140    170.5338    172.2131
## Isa_120301     172.3006      0.0000    178.9638    167.7560    175.3762
## Isa_120308     181.9140    178.9638      0.0000    169.8122    187.4563
## Isa_120320     170.5338    167.7560    169.8122      0.0000    163.9084
## Isa_120321     172.2131    175.3762    187.4563    163.9084      0.0000
```

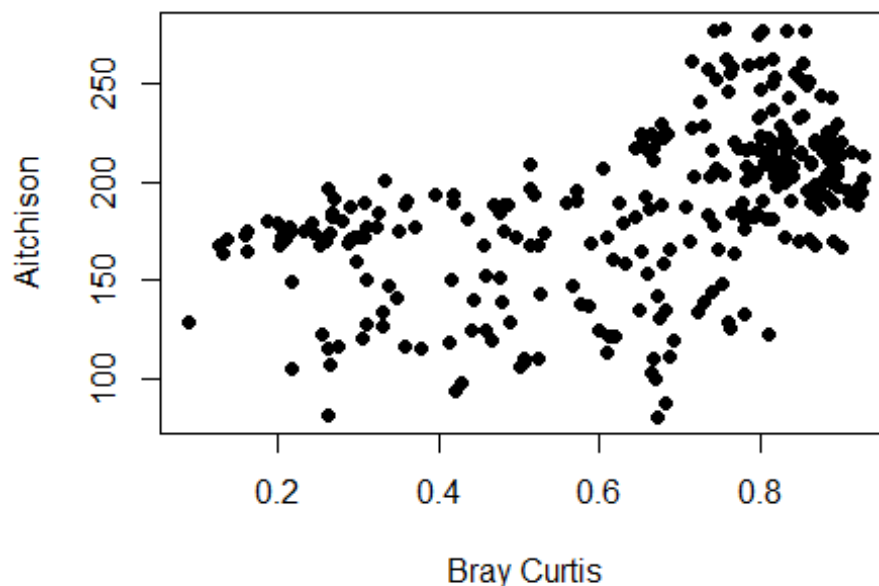
Let's compare the distance matrices:

```
identical(rownames(as.matrix(otu.tab.simple.ss.nozero.bray)),rownames(as.matrix(otu.tab.simple.gbm.clr.euclidean)))

## [1] TRUE
```

Generate a simple x-y plot, and fit the linear model (i.e. the regression)

```
plot(otu.tab.simple.ss.nozero.bray, otu.tab.simple.gbm.clr.euclidean, pch=19, xlab="Bray Curtis", ylab="Aitchison")
```



```
#lm<-lm(otu.tab.simple.gbm.clr.euclidean~otu.tab.simple.ss.nozero.bray)
#abline(lm, col="red")
```

The correlation between distance matrices is tested with a Mantel test.

```
mantel(otu.tab.simple.ss.nozero.bray, otu.tab.simple.gbm.clr.euclidean)

##
## Mantel statistic based on Pearson's product-moment correlation
##
## Call:
## mantel(xdis = otu.tab.simple.ss.nozero.bray, ydis = otu.tab.simple.gbm.clr.euclidean)
##
## Mantel statistic r: 0.5112
##      Significance: 0.001
##
## Upper quantiles of permutations (null model):
##      90%      95%  97.5%      99%
## 0.0722 0.1008 0.1386 0.1941
## Permutation: free
## Number of permutations: 999
```

Phew That was Part I. Now before you have a break save the data so it can be loaded if you want to use some of the same data.

```
save.image("AB332_lab_I.RData")
```